

Hi! PARIS Summer School 2023

Tutorial 6A

Scoring Strategically: Application to Finance

Johan Hombert (HEC Paris)

July 11, 2024, 14:00-17:30

Slides and data @ <https://johanhombert.github.io/fintech>

Road map

What is finance?

Business simulation

What is finance?

- Alice just graduated. She has a business idea with setup cost 100 K€ but no savings
 - Bob has 100 K€ of savings but no business idea
 - Without finance: Bob keeps his savings under his mattress. Alice does not start her business
 - Implications:
 1. Miss on good economic projects
 2. Savings do not earn returns
- ⇒ Inefficient allocation of resources

What is finance?

- With finance: Bob lends 100 K€ to Alice, who can launch her company
- Implications:
 1. Resources are allocated to good ideas
 2. Both investors (Bob) and borrowers (Alice) are better off
 3. Finance is a key input to economic development




Fig.: Countries with more developed credits market are richer
(caveat: correlation vs. causality)

Real-world financial arrangements

- Debt financing: Bob lends to Alice (fixed repayment + interests)
Equity financing: Bob takes a stake in Alice's business (dividends)
- Bob may invest in Alice's business directly or through a financial intermediary (bank, fund)
- Examples

	Debt	Equity
Direct	Friends and family Bond market	Angel investors Stock market
Intermediaries	Bank Debt mutual fund	Equity mutual fund Venture capital Private equity

The fundamental problem of finance

- Investors must assess if business ideas are good
- If Bob lends to Alice and Alice's project is worthless, then resources are wasted and would have better been kept under the mattress or lent to someone else

The fundamental problem of finance

A Decade After the Global Financial Crisis, Spanish Ghost Towns Remain

An estimated 3.4 million homes are currently unoccupied in Spain thanks to the country's great housing bust.



The fundamental problem of finance

- Investors must assess if business ideas are good

⇒ A prediction problem

How do investors do prediction in practice?



How do investors do prediction in practice?



BE A FINALIST OF HEC SEED PITCH COMPETITION, AND GET THE OPPORTUNITY TO
COMPETE AT THE AX-HEC ALUMNI COLLOQUIUM

*"BUSINESS
COMPETITIVENESS: OPENING
UP THE FIELDS OF VISION"*

ON NOVEMBER 12TH



ONLINE PITCH COMPETITION

How do investors do prediction in practice?



Use data

Use data to predict default: a case study

- “On the Rise of FinTechs: Credit Scoring Using Digital Footprints,” 2019, Berg, Burg, Gombovic and Puri, *Review of Financial Studies* [\[pdf\]](#)
-
- E-commerce company that did A/B testing
 - ▶ Randomize customers checking an item on the website
 - ▶ Treatment: offer option to pay within 15 days of receiving the purchased item
 - ▶ Control: option not offered
 - ▶ Same price of the item for both
- Impact
 - ▶ Control group: probability of buying the item = 45%
 - ▶ Treated group: probability of buying the item = 85%
- What should the management of company do?

Use data to predict default: a case study

- Tradeoff in extending payment facility: higher sales vs default risk
- ⇒ Estimate default risk and extend payment facility if estimated default probability is low
- Phase 1: credit score purchased from a credit bureau
 - ▶ Score based on credit history, payment history, sociodemographics
- Phase 2: the company realizes it has proprietary data on customers
 - ▶ Digital footprints: OS, email, log-in information, etc.
 - ▶ Does this data improve default prediction?

Scoring model

- Logistic regression to predict default
- Dependent variable: =1 if default
- Predictive variables

(1) Credit bureau score

(2) Digital footprints


(3) Credit bureau score + digital footprints

Default regressions (scorable customers)

Variables	(1) Credit bureau bureau score		(2) Digital footprint		(3) Credit bureau score & digital footprint	
	Coef.	z-stat	Coef.	z-stat	Coef.	z-stat
Credit bureau score	-0.17***	(-7.89)			-0.15***	(-6.67)
Device type & operating system ^d						
Desktop/Windows			Baseline		Baseline	
Desktop/Macintosh			-0.07	(-0.53)	-0.13	(-1.03)
Tablet/Android			0.29***	(3.19)	0.29***	(3.06)
Tablet/iOS			0.08	(1.05)	0.08	(0.97)
Mobile/Android			1.05***	(17.25)	0.95***	(15.34)
Mobile/iOS			0.72***	(9.07)	0.57***	(6.73)
E-mail Host ^a						
Gmx (partly paid)			Baseline		Baseline	
Web (partly paid)			0.00	(0.00)	-0.02	(-0.22)
T-Online (affluent customers)			-0.40***	(-3.90)	-0.35***	(-3.35)
Gmail (free)			0.34***	(3.81)	0.29***	(3.09)
Yahoo (free, older service)			0.75***	(9.19)	0.72***	(8.98)
Hotmail (free, older service)			0.35***	(3.70)	0.28***	(2.72)
Channel						
Paid			Baseline		Baseline	
Affiliate			-0.49***	(-5.35)	-0.54***	(-5.58)
Direct			-0.27***	(-4.25)	-0.28***	(-4.44)
Organic			-0.15*	(-1.79)	-0.15*	(-1.74)
Other			-0.47***	(-4.50)	-0.48***	(-4.36)
Checkout time						
Evening (6 p.m.-midnight)			Baseline		Baseline	
Morning (6 a.m.-noon)			0.28***	(4.50)	0.28***	(4.60)
Afternoon (noon-6 p.m.)			0.08	(1.42)	0.08	(1.47)
Night (midnight-6 a.m.)			0.79***	(7.73)	0.75***	(7.09)
Do-not-track setting			-0.02	(-0.25)	-0.07	(-0.91)
Name in e-mail			-0.28***	(-5.67)	-0.29***	(-5.70)
Number in e-mail			0.26***	(4.50)	0.23***	(3.91)
Is lowercase			0.76***	(13.10)	0.74***	(13.20)
E-mail error			1.66***	(20.00)	1.67***	(20.36)
Constant	12.42***	(5.76)	-4.92***	(-62.87)	9.97***	(4.48)
Control for Age, Gender, Item category, Loan amount, and month and region fixed effects	No		No		No	
Observations	254,819		254,819		254,819	
Pseudo R ²	.0244		.0524		.0717	
AUC	0.683		0.696		0.736	
(SE)	(0.006)		(0.006)		(0.005)	
Difference to AUC=50%	0.183***		0.196***		0.236***	
Difference AUC to (1)			0.013*		0.053***	

Impact on default rate

- The new credit scoring model is put in production in October 2015
- Default rate divided by 3



Road map

What is finance?

Business simulation

Business simulation

- You run a digital bank that gives loans to individuals
 - Principal amount: 10,000 euros given to borrower now, repaid by borrower in one year
 - Interest rate: i (in %) paid upfront
 - Default risk: the borrower defaults (=does not repay the principal) with some probability p
- Your cash flow is



⇒ Your **expected profit** for a given default probability:

$$-10000 + 10000 i + (1 - p)10000 + p \cdot 0 = (i - p)10000$$

Example

- Default probability: $p = 10\%$

Interest rate: $i = 6\%$

Expected profit: $(0.06 - 0.10) \times 10,000 = \text{loss of } 400 \text{ €}$

- Default probability: $p = 10\%$

Interest rate: $i = 12\%$

Expected profit: $(0.12 - 0.10) \times 10,000 = \text{gain of } 200 \text{ €}$

Loan offers

- You receive 100,000 loan applications
- Each loan applicant has a different default probability, which you don't know \Rightarrow You must estimate it from data (more on this later)
- You decide the interest rate i you offer to each loan applicant
- You are in competition with two other lenders (=two other teams), who also make loan offers to the same pool of applicants
- Loan applicants prefer a lower interest rate but have an intrinsic preference for one of the lender:
 - Let i_{k1} , i_{k2} , i_{k3} be the offers to applicant k from the three lenders
 - 1/3 of applicants have a preference for **lender 1** and choose the cheapest among $i_{k1} - 0.02$, i_{k2} , i_{k3}
 - 1/3 of applicants have a preference for **lender 2** and choose the cheapest among i_{k1} , $i_{k2} - 0.02$, i_{k3}
 - 1/3 of applicants have a preference for **lender 3** and choose the cheapest among i_{k1} , i_{k2} , $i_{k3} - 0.02$
 - Lenders don't know the preference of each applicant

Example

- Lender 1 offers 10%

Lender 2 offers 11.5%

Lender 3 offers 13%

- If applicant has a preference for lender 1 \Rightarrow chooses lender 1

If applicant has a preference for lender 2 \Rightarrow chooses lender 2

If applicant has a preference for lender 3 \Rightarrow chooses lender 1

Example

- Lender 1 offers 10%

Lender 2 offers 11.5%

Lender 3 offers 13%

- If applicant has a preference for lender 1 \Rightarrow chooses lender 1

If applicant has a preference for lender 2 \Rightarrow chooses lender 2

If applicant has a preference for lender 3 \Rightarrow chooses lender 1

- If the applicant's default probability is 11%, expected profit is

Lender 1: $\frac{2}{3} \times (0.10 - 0.11) \times 10,000 = \text{loss of } 66.67 \text{ €}$

Lender 2: $\frac{1}{3} \times (0.115 - 0.11) \times 10,000 = \text{gain of } 16.67 \text{ €}$

Lender 3: no gain no loss

Profit

- The goal is to maximize profit

$$\sum_{k=0}^{100,000} 1\{\text{Applicant } k \text{ takes your offer}\} \times (i_k - 1\{k \text{ defaults}\}) \times 10000$$

- The key is to estimate the default probability accurately and set the interest rate accordingly

Data

- `NewApplications_LenderX.csv` contains the 100,000 loan applications
- Lenders have **partially overlapping information** to predict default
- All three lenders have data
 - id: loan application identifier
 - sex: 1=male, 0=female
 - marital: 1=married, 0=other
 - employment: employment status (four categories)
 - income: annual income in euro (top coded at 1M euros)
- Lender $X = 1, 2, 3$ has data
 - digitalX: digital score from 0 to 1 as measured by lender X's proprietary algorithm

Data

- **PastLoans.csv** contains data on past loans with the following information
 - All the above variables, including digital1, digital2 and digital3 for all three lenders
 - default: 1=the borrower defaulted on the loan, 0=the loan was repaid
- You should use this data set to train a default probability model

Recap of data sets

	Lender 1		Lender 2		Lender 3	
	PastLoans.csv	NewApplications.csv	PastLoans.csv	NewApplications.csv	PastLoans.csv	NewApplications.csv
sex, marital, employment, income	✓	✓	✓	✓	✓	✓
digital1	✓	✓	✓		✓	
digital2	✓		✓	✓	✓	
digital3	✓		✓		✓	✓
default	✓		✓		✓	

Offers

- After you have estimated a default prediction model, you make loan offers to the 100,000 applicants
- Objective: maximize profit
- The maximum allowed interest rate is 25%
- You can choose not to make offers to some loan applicants
- Input the offers in a csv file with two columns
 - ▶ **id**: from the original dataset, from 0 to 99,999
 - ▶ **rate**: your interest rate inputted as a decimal number between 0 and 0.25 (input 0.12 for an interest rate of 12%). Leave the cell empty if you do not make an offer to a given applicant
- Name the csv file **teamN.csv** where N is your team number (between 11 and 35) and email it to hombert@hec.fr

How to set the interest rate?

- You should set the interest rate ABOVE the estimated probability of default because of the **winner's curse**
 - Each lender has different information \Rightarrow different estimate of default proba: p_1, p_2, p_3
 - If prediction models are unbiased, estimates are centered around the true default proba \Rightarrow the true proba will typically lie between $\min\{p_1, p_2, p_3\}$ and $\max\{p_1, p_2, p_3\}$
 - Suppose each lender offers interest rate = own estimate
 - The winning offer is the lowest one, which is below the true default proba \Rightarrow The winner always makes a loss!

How to avoid the winner's curve?

- Apply a margin of safety: interest rate $>$ estimated default proba
- In practice, the margin of safety can be calibrated by experimentation and back testing

Winner's curse in real estate investing



Zillow: Machine learning and data disrupt real estate

Learn how big data and the Zillow Zestimate changed and disrupted real estate. It's an important case study on the power of machine learning models and digital innovation.



Written by **Michael Krigsman**, Contributor
Posted in Beyond IT Failure on **July 30, 2017**

Interview with Zillow's Chief Analytics Officer Stan Humphries in 2017

ZD: How accurate is the Zestimate?

S.H.: Our models are trained such that half of the Earth will be positive and half will be negative; meaning that on any given day, half of [all] homes are going to transact above the Zestimate value and half are going to transact below.

Winner's curse in real estate investing



Zillow: Machine learning and data disrupt real estate

Learn how big data and the Zillow Zestimate changed and disrupted real estate. It's an important case study on the power of machine learning models and digital innovation.



Written by **Michael Krigsman**, Contributor
Posted in Beyond IT Failure on **July 30, 2017**

Interview with Zillow's Chief Analytics Officer Stan Humphries in 2017

ZD: How accurate is the Zestimate?

S.H.: Our models are trained such that half of the Earth will be positive and half will be negative; meaning that on any given day, half of [all] homes are going to transact above the Zestimate value and half are going to transact below.

CHRIS STOKEL-WALKER

BUSINESS NOV 11, 2021 8:00 AM

WIRED

Why Zillow Couldn't Make Algorithmic House Pricing Work

A few Nobel prizes

- The winner's curse is also called adverse selection or the lemon's problem
- Its study has been awarded several Nobel prizes



George A. Akerlof
(2001)



A. Michael Spence
(2001)



Joseph E. Stiglitz
(2001)



Paul R. Milgrom
(2020)



Robert B. Wilson
(2020)